



# Taking the Edge off with Espresso

Scale, Reliability and Programmability for Global Internet Peering

**KK Yap**, Murtaza Motiwala, Jeremy Rahe, Steve Padgett, Matthew Holliman, Gary Baldus, Marcus Hines, Taeun Kim, Ashok Narayanan, Ankur Jain, Victor Lin, Colin Rice, Brian Rogan, Arjun Singh, Bert Tanaka, Manish Verma, Puneet Sood, Mukarram Tariq, Matt Tierney, Dzevad Trumic, Vytautas Valancius, Calvin Ying, Mahesh Kallahalla, Bikash Koley, Amin Vahdat and many others.

# Problem Statement

Egress Terabits/sec of traffic to our Internet peers

- High-def video, cloud traffic, etc.

# Problem Statement

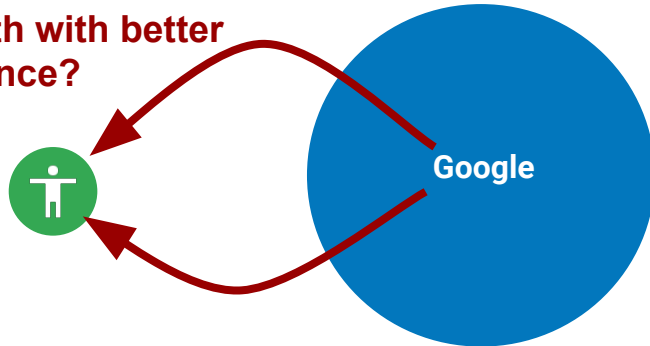
Egress Terabits/sec of traffic to our Internet peers

- High-def video, cloud traffic, etc.

## 1. Optimize traffic per-customer and per-application

- e.g., optimal video quality, or differentiated service for cloud
- **Problem: Constrained by BGP shortest path and lack of application awareness**

**Alternate path with better  
user experience?**



# Problem Statement


Egress Terabits/sec of traffic to our Internet peers

- High-def video, cloud traffic, etc.

## 2. Deliver new features quickly

- Problem: router-vendor feature cycles and qualification take many years

### Novel L2 VPN?



# Espresso: Google's SDN Peering Edge

Our previous experience with SDN

- B4 [SIGCOMM 2013] and Jupiter [SIGCOMM 2015]
- Enable flexible traffic engineering
- Increase feature velocity


SDN is only suited for walled gardens?

*Peering edge requires interoperability with heterogeneous peers.*

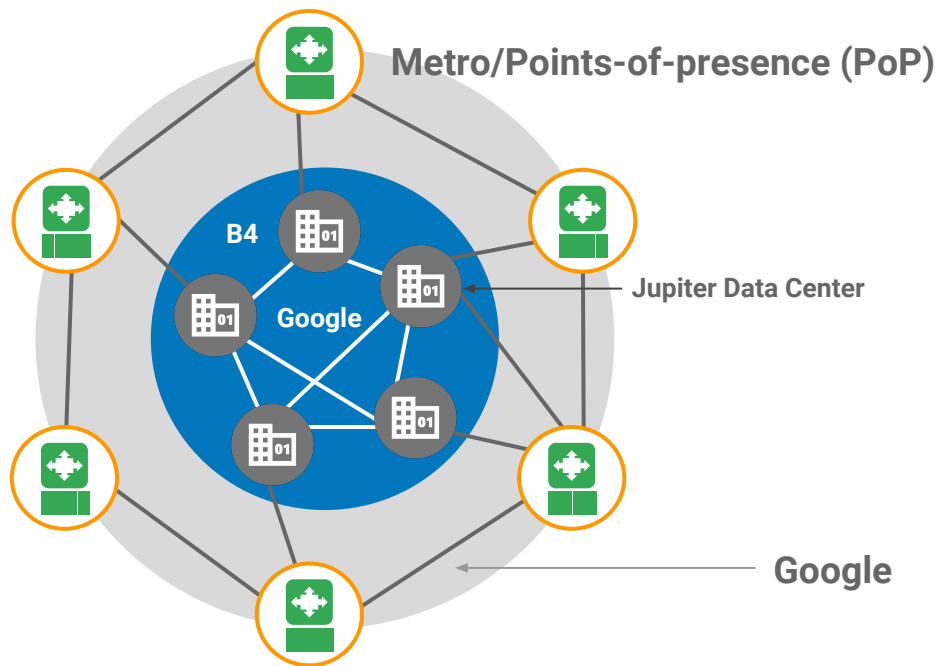
# Agenda

- Problem Statement
- Espresso in Context
- Design Principles
- Architecture Overview
- Results
- Conclusion

# Espresso in Context

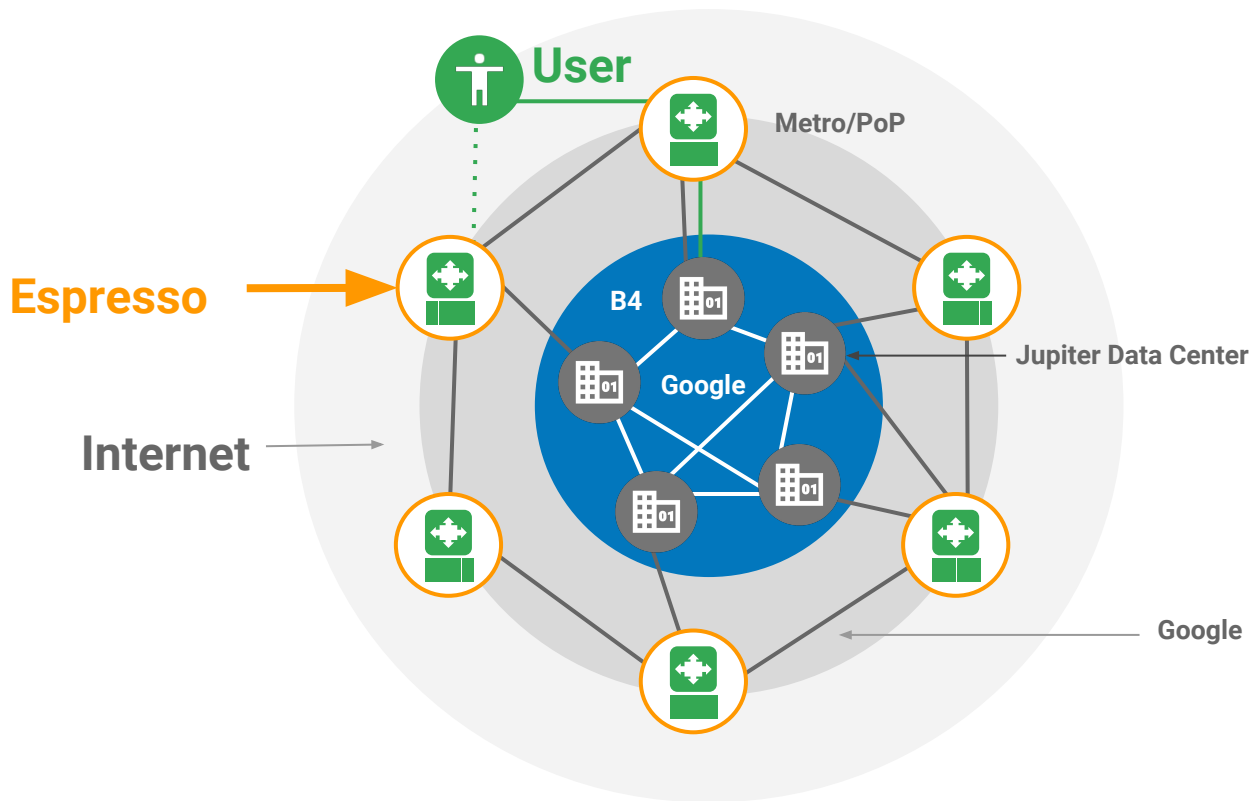


# Espresso in Context

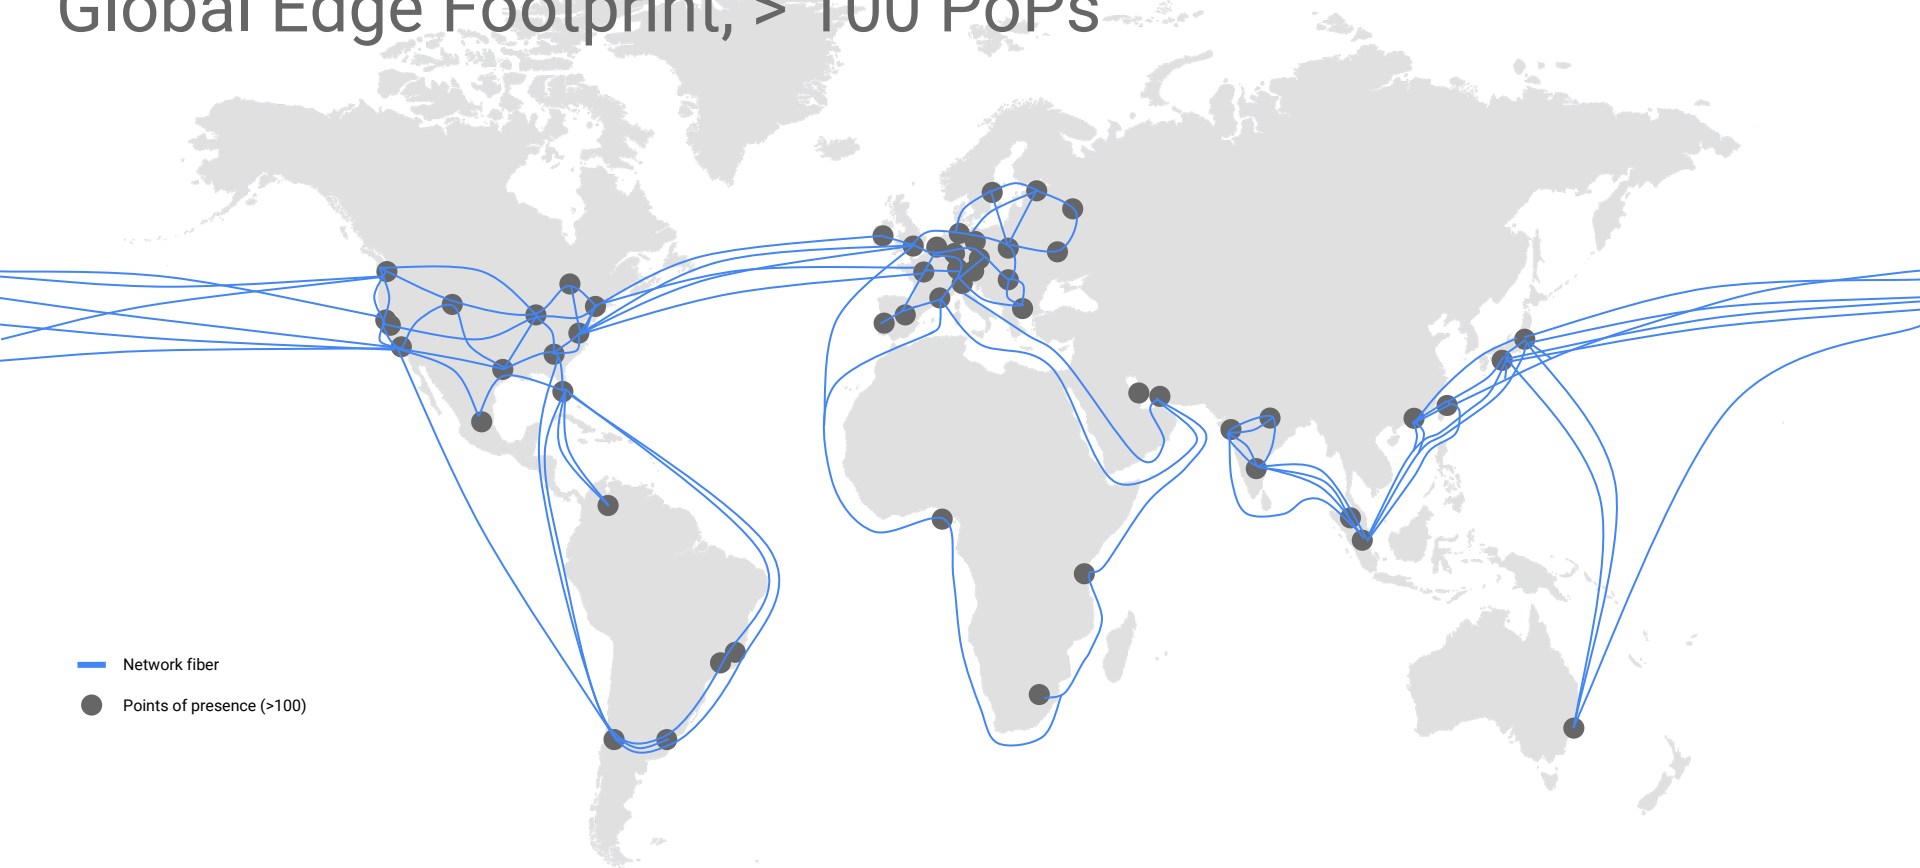




# Espresso in Context



# Global Edge Footprint, > 100 PoPs



# Agenda

- Problem Statement
- Espresso in Context
- Design Principles
- Architecture Overview
- Results
- Conclusion

# Espresso's Design Principles

## 1. Hierarchical control plane

- Global optimization while local control plane provide fast reaction.

## 2. Fail static

- Local control plane continues to function without global controller failure.

## 3. Software programmability

- Externalize features into software to exploit commodity servers for scale.

## 4. Testability

## 5. Manageability

# Espresso's Design Principles

## 1. Hierarchical control plane

- Global optimization while local control plane provide fast reaction.

## 2. Fail static

- Local control plane continues to function without global controller failure.

## 3. Software programmability


- Externalize features into software to exploit commodity servers for scale.

## 4. Testability

## 5. Manageability

# Architecture: Externalizing BGP


## Traditional Peering Router



eBGP Peering




Host Servers  
in Metro




## Espresso

Peering Fabric

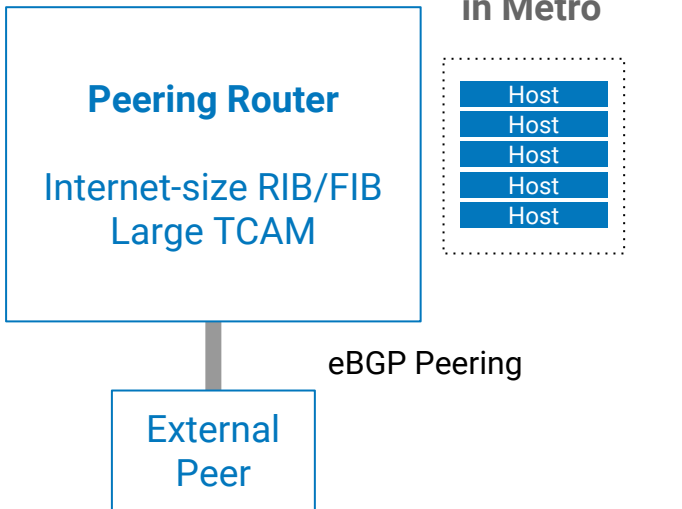


Host Servers  
in Metro

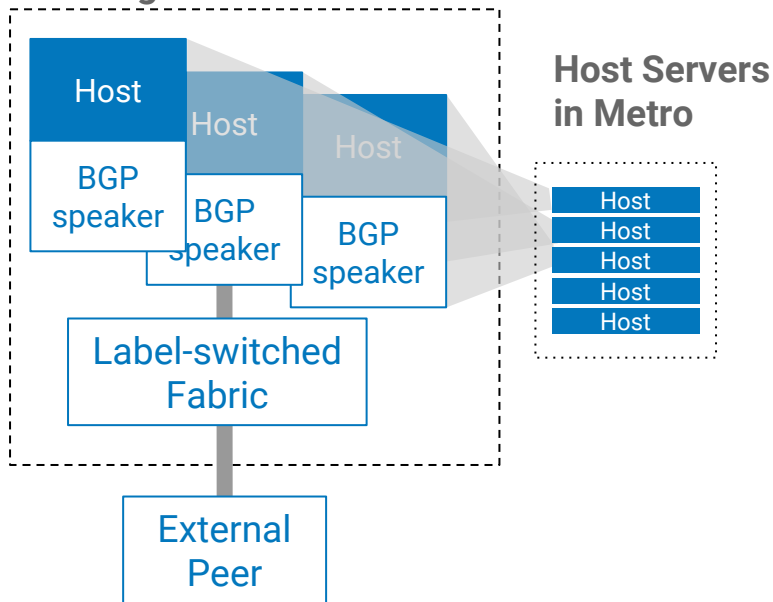


# Architecture: Reliability and Scale of BGP

## Traditional Peering Router

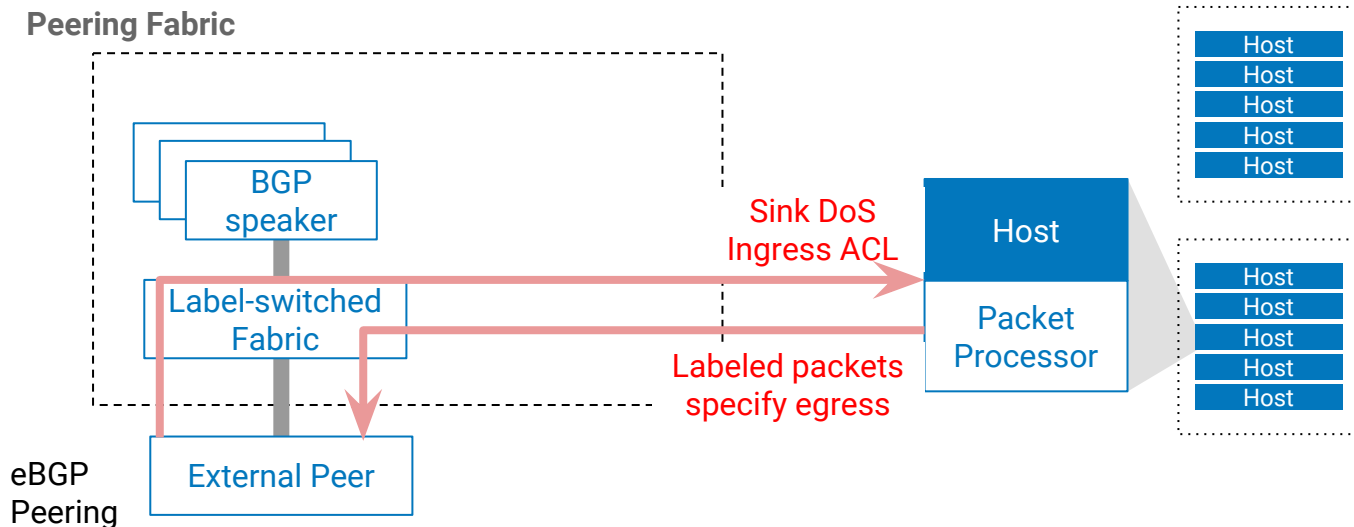


## Espresso Peering Fabric



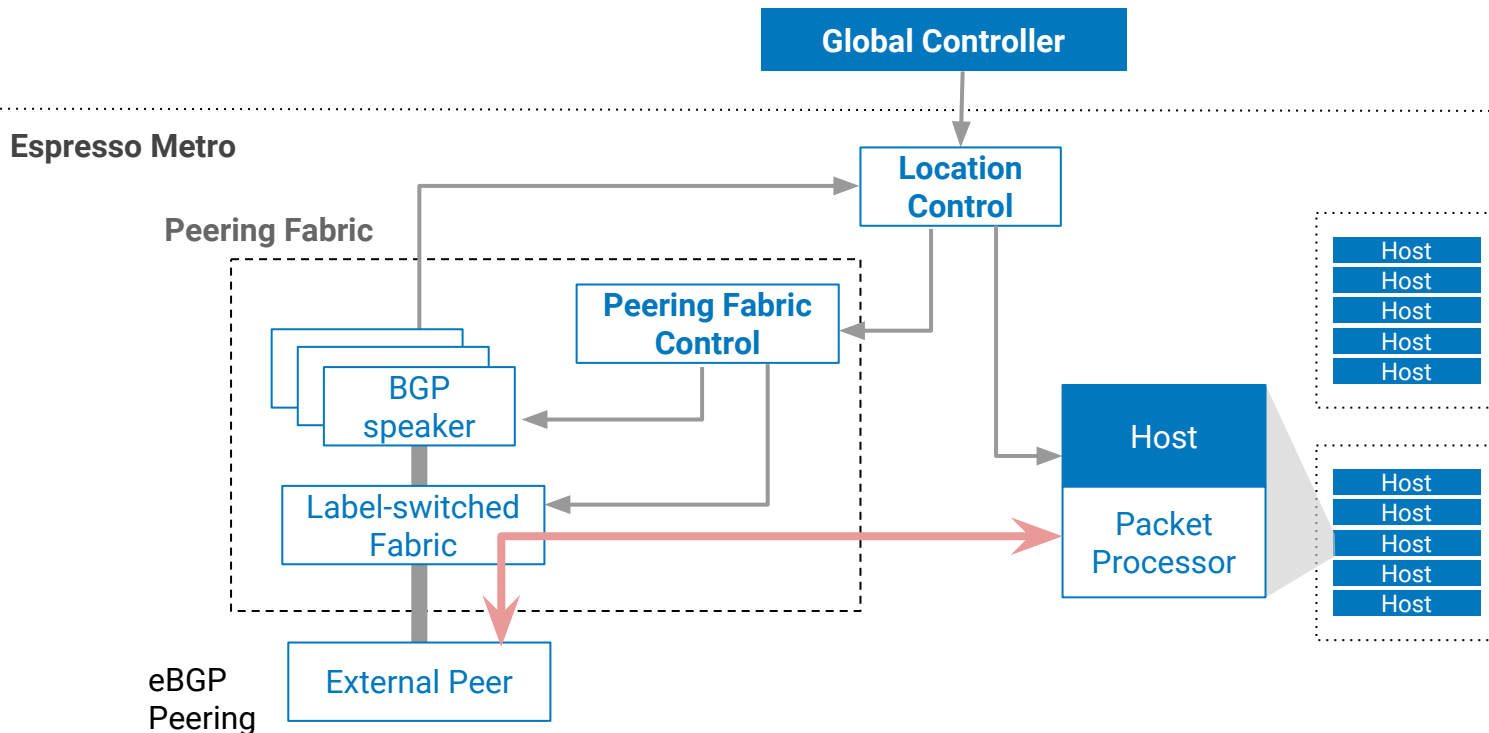
# Architecture: Externalize Packet Processing

Host-based packet processor allows flexible packet processing, including ACL and handling of DoS.

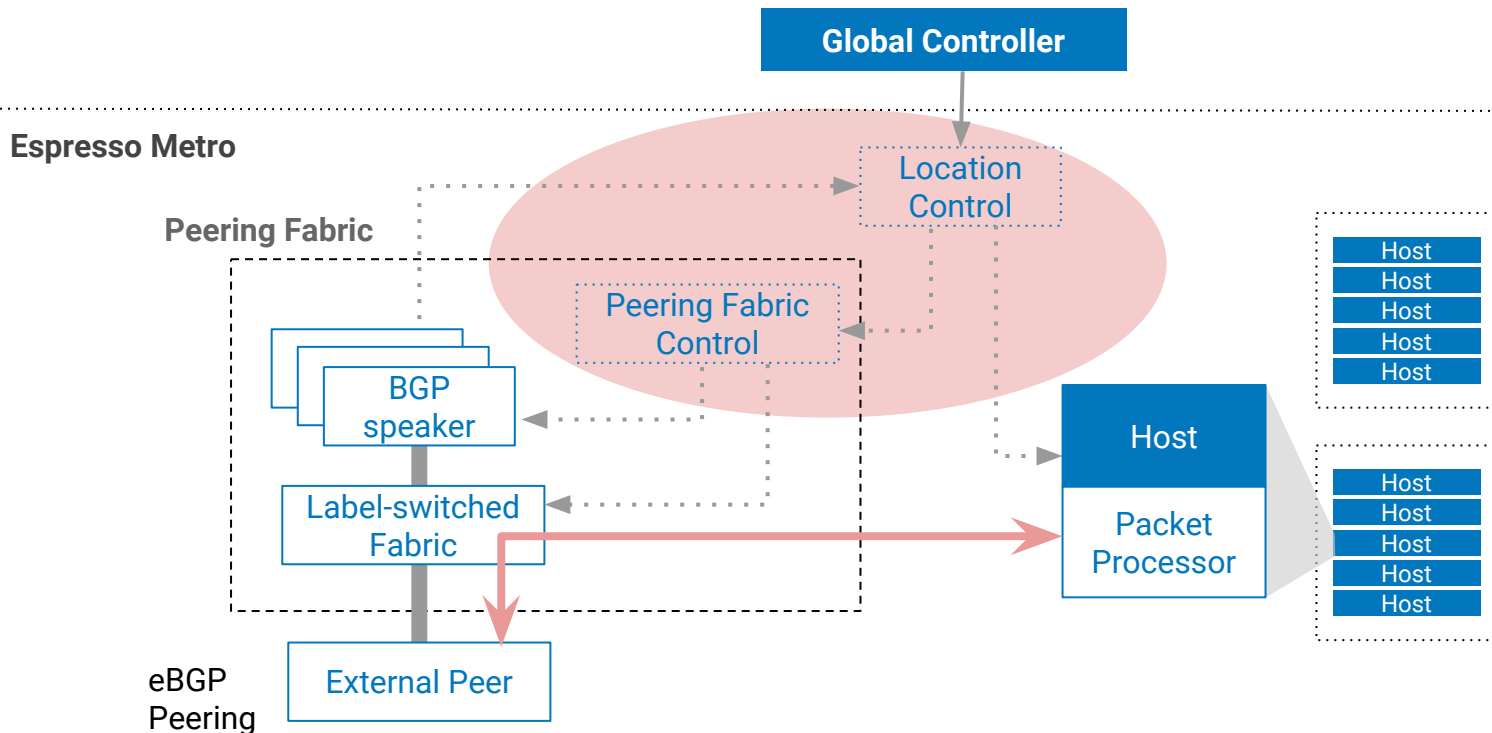




# Architecture: Hierarchical Control

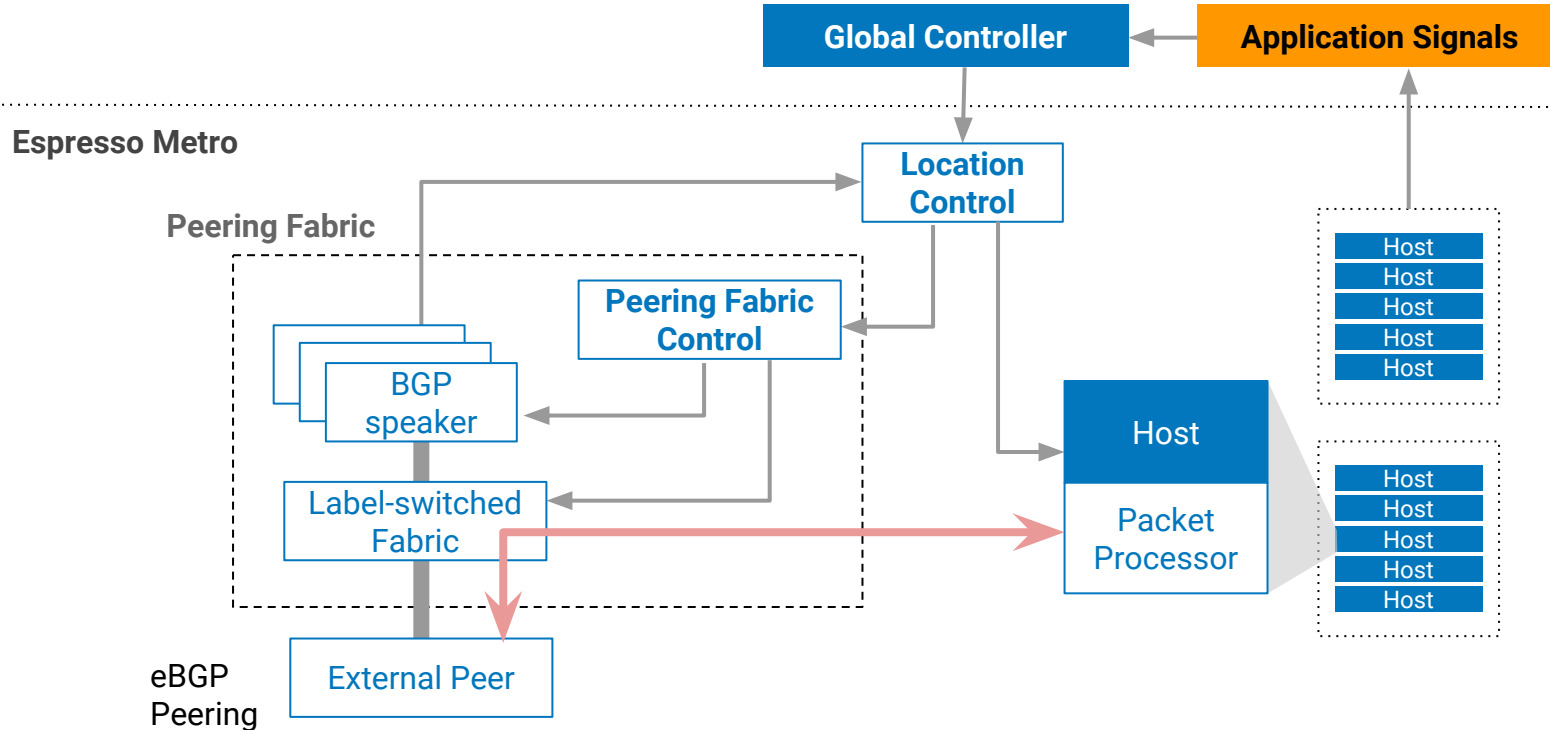


# Architecture: Fail Static




*Hierarchical control plane*  
*Fail static*  
**Software programmability**

# Architecture: Application Aware Routing



# Using User's Best Path, not BGP's



- Serve **13% more traffic** than BGP best path in application aware manner.
- Helps capacity-constrained ISPs by overflowing demand to alternate paths within local metro and **also via remote metros.**

# Improvements in End User Experience

<b>Client ISP</b>	<b>Change in mean time between rebuffers (MTBR)</b>	<b>Change in Mean Goodput</b>
A	10 → 20 min	2.25 → 4.5 Mbps
B	4.6 → 12.5 min	2.75 → 4.9 Mbps
C	14 → 19 min	3.2 → 4.2 Mbps

**Provide significant improvements to end-user experience.**

# Release Velocity

Component	Average Velocity (days)
Local Controller	11.2
BGP speaker	12.6
Peering Fabric Controller	15.8

**> 50× more frequently than with traditional peering routers.**

**Novel L2 VPN delivered 6× faster via incremental rollout.**

# Conclusion

~~SDN is only suited for walled gardens.~~

Espresso demonstrates that

- traditional peering architecture can evolve to exploit SDN
- SDN's value is in flexibility and feature velocity

# Conclusion



Local view

Connectivity based optimization

Slow evolution

Costly

Global view

Application signals-based optimization

Rapid deploy-and-iterate

75% Cheaper